THE
PHANTASM
CRASHCOURSE

A Complete Introduction to the PHANTASM Grammar
for WebAssembly Programmers

IMPORTANT: *The PHANTASM Project is in its absolute infancy. Anything published at this stage is only made available as a preview. The project uses GitHub Issues as a general forum, so feel free to start a discussion there, if you have any questions.*

The *Portable, Hardened, Asynchronous, Natively Typed, Abstract Stack Machine* (*PHANTASM*) is a novel web assembler that allows you to author WebAssembly modules using a nice, modern syntax, rather than repurposing the Text Format.

This crashcourse summarizes the entire PHANTASM grammar in terms that anyone familiar with WAT should understand. The project wiki contains a growing collection of articles that expand on this document, covering finer details that are beyond the scope of an introductory tutorial.

Reading this document should be enough to begin programming with PHANTASM (assuming you already have a working knowledge of WAT and WebAssembly), and the Wiki can be referenced for anything that remains unclear.

Source Files

PHANTASM source files use the `phantasm` file extension, or the alias `ph`. Longer names like `main.phantasm` are preferable, but names like `main.ph` can be used whenever an eight-character extension would be too unwieldily (and disambiguation from Perl header files is not an issue).

PHANTASM binaries are just WebAssembly binaries, and therefore use the `wasm` extension, though the filepath (including its extension) is preserved in the *Module Name* subsection of the *Custom Name Section*.

Character Set

Source files use a safe subset of ASCII. The legal character set includes all of the ASCII printables (including the Space Character), as well as the Newline Character. It does not include the Tab Character (please use soft-tabs).

Note: String literals support escape sequences that make it easy to express Unicode characters using the ASCII subset available in a PHANTASM source file (this is documented below).

Special Characters

Like WAT, PHANTASM uses the longest-match rule with the same set of special characters (whitespace, commas, semicolons, parens, brackets and braces). However (outside of strings and comments), PHANTASM does not actually use semicolons, parens, brackets or braces, and tabs are illegal.

The three special characters that are used in code (space, newline and comma) are the characters that structure the code, while the regular characters are used to spell the proper tokens.

Significant Whitespace

PHANTASM uses significant whitespace and the Comma Character to delimit tokens, logical lines and blocks.

A Newline Character always terminates the logical line, unless the ellipsis-operator (see below) is used to explicitly prevent termination.

The Space Character is used to delimit individual tokens, and to indent blocks. Indentation must contain *exactly* four spaces per level, always.

The Comma Character is used to delimit sequences, and to compound repetitious constructs to minimize redundancy and improve readability in various contexts, each individually described in the relevant sections of the tutorial.

Note: The `end` pseudo-instruction is not used in PHANTASM.

Special Tokens

The longest match rule applies to all PHANTASM tokens, except comments (inline and multiline) and string literals, which are all handled as special-cases.

Inline-comments start with a semicolon (`;`), then run to the end of the line, while multiline-comments start and end on a double-semicolon (`;;`). For example:

```
define $sum of type $binop

    ;; This is multiline-comment. It is being used as
    a docstring for a function-definition. ;;

    get 0, get 1, add i32 ; this is an inline comment
```

String Literals

Like WAT, PHANTASM uses the Quote Character (`"`) to delimit string literals:

```
"Hello, World!"
```

Unlike WAT, neither slash character has any special meaning inside PHANTASM string literals. DOS-paths can be written normally:

```
"c:\windows\command"
```

Curly braces are special in a PHANTASM string literal. They are used to wrap one or more space-separated *escape-expressions*, which together form an *escape-sequence*. For example, this string ends with a grinning-face emoji:

```
"nice one! {1F600}"
```

And this string ends by placing a grinning-face emoji on a newline:

```
"nice one!{n 1F600}"
```

To include actual braces in a string literal, they must be doubled-up (or expressed with an escape-expression). For example, the following string is valid CSS:

```
"body {{margin: 0}}"
```

Each escape-expression (within a given escape-sequence) expresses exactly one Unicode character, either using its Unicode codepoint, expressed in uppercase hexadecimal, or using a lowercase name (or a shorter alias), defined by the assembler.

See the wiki article *String Literals* for more information, including all of the currently defined character-names, and a note on adding new ones.

Regular Tokens

Every regular token (every token that is not a string or a comment) follows the longest-match rule, and is then classified as belonging to one of five token types: *keyword*, *mnemonic, operator, identifier* or *number literal*.

Keywords

Keywords are used to begin statements, and to begin continuations within statements. They are the most important tokens in the language, and provide the various statements of the language with a consistent structure.

Keywords are always simple English words, spelled in lowercase. The spelling is not enough to disambiguate keywords from mnemonics, so regular tokens are classified as keywords if they are in the following list:

```
define    import    export    with      from      thus      via
of        to        at        in        as        sop
```

Note: A policy on reserved words would be redundant.

Mnemonics

Mnemonics are simple, little names that the language assigns to its instructions, components, types *et cetera*. This includes the names of packed types (like `s8` and `utf8`), and the qualifiers used to qualify other mnemonics (as in `pointer table`, `shared memory` *et cetera*).

Note: PHANTASM spells the WAT reftypes `funcref` and `externref` as `pointer` and `proxy` respectively.

Mnemonics are lexically similar to keywords, except that mnemonics may also contain digits (but will not begin with one). All of the mnemonics are listed below:

```
f32       f64       utf8      void
i8        s8        u8        shared
i16       s16       u16       segment
i32       s32       u32       zero
i64       s64       u64       null
pointer   proxy     type      bank
variable  constant  register  start
function  memory    table     mixed
global    local     left      right
equal     more      less      nop
call      invoke    return    unreachable
block     loop      branch    else
jump      fork      exit      select
get       set       put       copy
```

```
load        store       drop        push
abs         neg         nearest     ceiling
floor       root        min         max
sign        wrap        cast        atomic
notify      wait        fence       swap
broker      is          not         nsa
clz         ctz         grow        fill
add         sub         mul         div
rem         and         or          xor
size
```

Operators

PHANTASM currently has two operators: the arrow-operator (`->`) and the ellipsis-operator (`...`).

Operators are spelled using non-alphanumeric characters.

The arrow-operator is used to map arities to results when defining the types and signatures of functions. For example:

```
import "sum" as $sum of i32, i32 → i32
```

The ellipsis-operator is used to continue a logical line across two or more lines of a source file. The operator prefixes each newline that should be ignored, and indents each continuation-line one level (relative to the opening line). For example:

```
import "pseudo.private.static.thing" from "/static/library.js" ...
    as function of i32, i32, i32, i32 → ...
    pointer, pointer, pointer
```

Note: Trailing whitespace after the ellipsis-operator is insignificant.

Identifiers

PHANTASM identifiers are identical to WAT identifiers (`$foo`, `$bar` *et cetera*).

Taking advantage of the fact that identifiers are always constant (they cannot be rebound), the assembler defers the resolution of identifiers until all of the indices (within each indexspace) have been assigned, and any identifiers bound to their respective indices.

Deferring identifier resolution permits identifiers to be referenced before they are bound (just like explicit indices). This convenience allows statements to be ordered in the most natural, readable way (which is not always possible in WAT code).

Note: PHANTASM uses the term *identity* to refer to an index, expressed either directly (with a number literal) or indirectly (with an identifier bound to the index). There are many constructs in the language that accept or require an identity.

Number Literals

PHANTASM number literals are similar to WAT, except that PHANTASM uses a hash (`#`) prefix (instead of `0x`) for hexadecimals, and PHANTASM has its own syntax for exponentiation (with its own semantics for hexadecimal exponentiation).

Any PHANTASM number literal (integer or float, decimal or hexadecimal) can be suffixed with the raise-operator (`\`) or lower-operator (`/`), followed by any natural number, which causes the value to be raised or lowered by that number of orders of magnitude. For example:

```
1\3      ; equal to 1,000
1/3      ; equal to 0.001
1.5\6    ; equal to 1,500,000
#FF\6    ; equal to 4,278,190,080, or #FF000000
#1.F/4   ; equal to #0.0001F
```

Note: Unlike WAT (and the C language-family generally), in PHANTASM the number after the operator is alway expressed using both the same base and the same notation as the number before the operator:

```
#FF\A     ; equal to #FF0000000000
#1.F/10   ; equal to #0.0000000000000001F
10\A      ; unrecognized token (`A` is not a decimal digit)
```

The floating-point constants (positive infinity, negative infinity and not-a-number), are expressed exactly like they are in JavaScript (`Infinity` or `+Infinity`, `-Infinity` and `NaN`).

The wiki article *Number Literals* provides more details, including how literals are evaluated and encoded, depending on the context (for example, `i32 1` and `f32 1` use the same literal, but the compiler encodes each one differently). The article also covers edgecases, like using floating point notation with an exponent to express an integer (which affects the range of integers that can be expressed).

We have already covered the basics of PHANTASM grammar (significant whitespace, special characters and the longest-match rule), as well as comments and the five regular token types (keywords, mnemonics, operators, identifiers and numbers).

The rest of the document describes the PHANTASM module grammar, from the top down.

The Module Grammar

Each PHANTASM file (implicitly) defines exactly one WebAssembly module.

At the top level, a module is simply an array of *statements*. Each statement occupies its own logical line (they are never indented, nor grouped together).

Each statement does one of three things to a *component*: A statement can *define*, *import* or *export* a component, and correspondingly, there are three types of statement (*define-statements*, *import-statements*, and *export-statements*).

The components are divided into two types: *system-components* (*registers*, *functions*, *memories* and *tables*) that can be defined, imported and exported, and *internal-components* (*function types*, *memory banks* and *table banks*) that can only be defined.

Statements

Of the three kinds of statement, define-statements have the simplest grammar:

define `component-definition`

Import statements have the most complicated grammar, as they have an optional part that defines the module-name string (that defaults to `"host"` when omitted):

import `"field"` [**from** `"module"`] **as** `component-specifier`

Export statements use the following grammar:

export `"field"` **as** `component-reference`

The three statement-grammars describe the entire language at the top level:

```
module : statement *
statement : define component-definition
          | import "field" [from "module"] as component-specifier
          | export "field" as component-reference
```

The three *component-descriptors* (*component-definitions*, *component-specifiers* and *component-references*) complete the statement grammar.

Note: The *component-definition* grammar (used by define-statements) is a pure superset of the *component-specifier* grammar (used by import-statements), while the *component-reference* grammar (used by export-statements) is naturally very simple (just one or two tokens). So, while there are seven kinds of component, and three different contexts, the various *component-descriptors* this requires all share a relatively simple and consistent syntax.

Component References

Of the three kinds of component-descriptor, the component-reference has the simplest grammar:

```
component identity
```

It is simply the component type, followed by its identity. For example:

```
register 1
function $sum
memory 0
table $opcodes
```

Memory-references and table-references can omit their identity, and zero will be inferred, so for example, the following memory-references are equivalent:

```
memory 0
memory
```

Component-references are used in export-statements. For example:

```
export "score" as register $score
export "RAM" as memory
```

While types and banks (function types, memory banks and table banks - see below) are (at least currently) all internal-components (they cannot be imported or exported), there are other places in the language where these types of components are referenced, and the same component-reference grammar is used. For example:

```
type $binop
memory bank 0
table bank $extensions
```

These internal component-references are used in various statements and instructions. For example:

```
import "helper" as $helper of type $binop
```

```
invoke type 5
copy memory bank $messages
```

Note: References to memory banks and table banks cannot omit their identities, like memories and tables can.

Component Specifiers

Component-specifiers are used in import-statements. As such, there are four kinds, one for each kind of system component (register, function, memory and table).

Register Specifiers

Register-specifiers prefix the valtype of the register with a register qualifier (either `variable` or `constant`), and can optionally append an identifier, which will be bound to the newly imported register:

```
register-qualifier valtype [identifier]
```

For example:

```
variable i32
constant pointer $pointer
```

Register-specifiers are used in import-statements. For example:

```
import "score" as variable i64
import "PI" from "/std/math" as constant f64 $PI
```

Function Specifiers

Function-specifiers use the `of` keyword, with two slightly different grammars, depending on whether an identifier is bound to the newly imported function or not. Unidentified function-specifiers use the following syntax:

```
function of type
```

Identified function-specifiers use this syntax:

```
identifier of type
```

The type of the specified function can be described with a *type-reference* (as described above, in the *Component References* section), or with a *type-expression* (described next, in the *Type Expressions* section). For example:

```
function of type $binop
$sum of i32, i32 → i32
```

Function-specifiers are used in import-statements. For example:

```
import "sum" from "mathlib" as $sum of type $binop
import "percent" as function of f32, f32 → f32
```

Type Expressions

Type-expressions express function types by mapping one list of reftypes (the arity) to another (the results), using the arrow-operator (`->`).

```
arity → results
```

The arity and results are expressed as comma-separated lists of reftypes, or `void` when the list is empty. For example:

```
i32, i32 → f64
pointer → i64
f32, proxy → void
```

Type-expressions are used in various statements and instructions. For example:

```
import "pointer.get" as $get_pointer of i32 → pointer

invoke i32, i32 → i32
```

Note: Function-definitions can use a *signature* to define their type (explained below).

Memory Specifiers

Memory-specifiers begin with an optional `shared` qualifier, followed by the component name `memory`. The component name is followed by an optional identifier (to be bound to the newly imported memory), which is in turn followed by the *limits* of the memory:

```
[shared] memory [identifier] limits
```

The limits-construct starts with the keyword `with`, followed by a number literal that defines the minimum length:

```
with number-literal
```

The previous grammar is enough for memories that do not define a maximum length. For those that do (noting that shared memories *must* define a maximum length), the previous grammar can be extended in one of two ways:

```
with number-literal to number-literal
```

```
with number-literal max
```

In the later case (where `max` is used), the maximum length is inferred to be the same as the minimum length), so the following two limits-constructs are equivalent:

```
with 32 to 32
with 32 max
```

Below are some examples of memory specifiers (each specifying its limits):

```
memory with 1

memory $store with #10 to #20

shared memory $current_track with 16 max
```

Memory-specifiers are used in import-statements. For example:

```
import "ram" from "sys" as memory with #10 to #20

import "audio" as shared memory $audio with 16 max
```

Table Specifiers

Table-specifiers use a grammar that is almost identical to memory-specifiers. However, in a table-specifier, the qualifier is required, and must be one of `pointer`, `proxy` or `mixed` (equivalent to the WAT reftypes `funcref`, `externref` and `anyref`, respectively). For example:

```
pointer table

proxy table

mixed table
```

As with memories, the component name is followed by an optional identifier (to be bound to the newly imported table), followed by the limits:

```
table-qualifier table [identifier] limits
```

Note: The minimum and maximum lengths (within a limits-construct) are expressed as the number of slots (not pages, as with memories) when used to specify a table.

Below are a few examples of table-specifiers:

```
proxy table with 256

pointer table $opcodes with #100 to #200

mixed table $extensions with 16 max
```

Table-specifiers are used in import-statements. For example:

```
import "opcodes" as pointer table $opcodes with #100 to #200

import "extensions" from "sys" as proxy table with 256 max
```

Component Definitions

Component-definitions are used in define-statements.

Along with component-references and component-specifiers (covered above), component-definitions make up the three component-descriptors.

Component-definitions reuse the grammar of the component-specifiers (used by import-statements), extending it as required. However, those grammatical extensions include the bodies of functions (with all of the instructions), as well as the *primers* that are used to populate memories, tables and banks (with all of the data and table elements they include), so there is still quite a lot of ground to cover.

Register Definitions

Register-definitions append an optional initializer to the grammar for register-specifiers:

```
register-qualifier valtype [identifier] [register-initializer]
```

The initializer is used to define the constant expression that the engine will use to initialize the register.

When the initializer is omitted, numtype registers are initialized to zero, while reftype registers are initialized to null (a constant expression containing the appropriate instruction is generated by the assembler automatically). For example, the following statement defines a writable register for a 32-bit integer, which will be initialized to zero (and identified as `$score`):

```
define variable i32 $score
```

When an initializer is provided, it can be expressed with a block of one or more instructions. For example:

```
define variable f64 $master.gain

    push f64 100
```

Single-line blocks can be inlined using the `thus` keyword. For example:

```
define variable f64 $master.gain thus push f64 100
```

Register-definitions also support sugar for blocks that simply push a constant (as in the example above). This grammar uses the `as` keyword, and allows numtype registers to be initialized using a number literal (without requiring a block):

```
define variable f64 $master.gain as 100
```

The `as` keyword can also be used to initialize function-registers, except here, the initial value is an identity (which is compiled to the `ref.func` instruction):

```
define variable pointer $foo as $bar
define constant pointer as 1
```

The `as` keyword cannot be used to initialize proxy-registers, as there is no analog of `ref.func` for proxies (externrefs) that the sugar could sensibly compile to.

Function Definitions

Function-definitions append a required initializer (the body of the function) to the grammar for function-specifiers, and replace the type-construct with a *signature-construct*, producing the following grammar for unidentified function-definitions:

function of `signature body`

The following variant is used for identified function-definitions:

`identifier` **of** `signature body`

The signature-construct extends the grammar of type-expressions, permitting (though never requiring) an identifier after each reftype in the arity-construct. For example:

```
i32 $x, i32 $y, i32 $z → i32
```

The signature grammar also permits the type of a given parameter to be inferred from the previous parameter (recursively). So, the previous example could be shortened to the following signature:

```
i32 $x, $y, $z → i32
```

The `body` construct begins with zero or more *local-directives*, followed by a block of instructions.

Statements, Directives & Commands

Each top-level construct in the grammar is a *statement*, a *directive* or a *command*.

Statements were described already (in the *Statement Grammar* section above). In short, they define a module at the top level, while directives and commands are always contained within blocks.

Directives must occupy their own line (like a statement) within the containing block. There are two kinds of directive: *local-directives* (described below) and *segment-directives* (explained later, in the *Memory Definitions* and *Table Definitions* sections).

Commands include the *instructions* that define the logic of functions, the *encodings* that populate memories and memory banks, and the *references* that populate tables and table banks. Commands can be grouped on the same line, using a comma-separator.

Local Directives

Local-directives begin with the local mnemonic, followed by the valtype of the local register, which is in turn followed by an optional identifier:

```
local valtype [identifier]
```

Below are a couple of simple examples:

```
local i32
local pointer $helper
```

Multiple local registers can be declared in a single directive. For example:

```
local i32 $x, i32 $y, i32 $z
```

As with parameters, the type can be inferred from the previous element. So, the last example could be shortened to this:

```
local i32 $x, $y, $z
```

A single local-directive can define different types of local register too. For example:

```
local i32 $a, $b, $c, i64 $x, $y, $z
```

As with parameters, when an identifier is omitted, the type is required:

```
local i32, i32, i32, i64, i64, i64
```

The Instructions

After any local-directives, the body of a function-definition contains one or more *instructions*.

Instructions are commands. As such, they can be grouped with other instructions, using a comma-separator. This includes the block-instructions (`block`, `loop`, `branch` and `else`), though block-instructions (obviously) end with a block of instructions, so naturally, nothing can be grouped after a block-instruction (only before one).

Note: To avoid the issues that CoffeeScript had with cryptic whitespace, the assembler enforces a rule that states that block-instructions cannot use inline blocks, and inline-blocks cannot use block-instructions. Otherwise, stuff like this becomes permissible:

```
define function of i32 → void thus branch of type 0
    call $foo
else thus call $bar
```

The individual instructions are documented in the following sections.

The WAT Mnemonics

The following five instructions all use single-token mnemonics that are copied directly from WAT:

```
unreachable
nop
return
else
drop
```

The Get, Set & Put Mnemonics

PHANTASM uses the mnemonics `get` and `set` for the WAT instructions `global.get`, `global.set`, `local.get`, `local.set`, `table.get` and `table.set`.

Both `get` and `set` take an optional scope-qualifier, which when present, must be one of `global`, `local` or `table`, defaulting to the lexical scope of the instruction, so `local` within functions, and `global` otherwise:

```
mnemonic [scope-qualifier] identity
```

Below are a few examples of the `get` and `set` mnemonics:

```
get $x
set local 6
get table $opcodes
```

PHANTASM uses the `put` mnemonic for the WAT `local.tee` instruction. Unlike `get` and `set`, `put` does not take a scope-qualifier:

```
put identity
```

The Push Mnemonic

PHANTASM has a `push` mnemonic that is used for all of the WAT instructions that push a value to the stack: `i32.const`, `i64.const`, `f32.const`, `f64.const`, `ref.func` and `ref.null`.

The `push` mnemonic uses the following grammar:

```
push [valtype] immediate
```

The valtype defaults to `i32`, and the immediate must be valid for the (given or implied) valtype. When the valtype is a numtype, the immediate must be an appropriate number literal. When the valtype is `pointer`, the immediate must be an identity or `null`, and when the valtype is `proxy`, the immediate must be `null`.

Below are examples of how the `push` mnemonic is used more generally:

```
push #30                 ; i32.const 0×30
push f64 0.5             ; f64.const 0.5
push pointer $helper     ; ref.func $helper
push pointer null        ; ref.null func
push proxy null          ; ref.null extern
```

Note: The WAT instruction `ref.is_null` is handled by the PHANTASM instruction `is null` (described below).

The Load & Store Mnemonics

PHANTASM uses the `load` and `store` mnemonics for all of the WAT load and store instructions, including those that load or store a *datatype*.

PHANTASM uses the term *datatype* to describe the types that can be written to, or read from, linear memory, that must be extended or truncated (to a proper numtype) in the process of moving them to or from the stack.

The `load` and `store` mnemonics use the following grammar:

```
mnemonic numtype [as datatype] [in identity] [at number-literal]
```

The `as` keyword, when present, prefixes the datatype (when it differs from the given numtype). The `in` keyword, when present, prefixes the identity of the memory (supporting multiple memories). The `at` keyword, when present, prefixes the offset of the load or store into memory.

For example:

```
load i32                    ; i32.load
store f64                   ; f64.store
store i64 as u32            ; i64.store32_u
load i64 as i8 at 3        ; i64.load8 offset=3
store i32 as s8 in $memory ; i32.store8_s $memory
load i64 as i8 in $RAM at 1 ; i64.load8 $RAM offset=1
```

Note: PHANTASM never uses signed or unsigned mnemonics (as WAT does). Instead, PHANTASM always expresses whether a given instruction operates on signed, unsigned or agnostic integers using corresponding signed, unsigned and agnostic integer-type names.

Note: The atomic instructions generally use the same grammar as the `load` and `store` instructions. This is fully documented below.

The Arithmetic & Logic Mnemonics

PHANTASM uses essentially the same grammar for all arithmetic operations:

```
mnemonic ~numtype
```

The valid set of numtypes depends on the individual instruction. Some instructions operate on both integers and floats, while others only work with integers or floats. Furthermore, of those that operate on integers, some require explicitly signed or unsigned integers, while others are type-agnostic.

The following list of mnemonics all use the above grammar, and have names that are copied directly from WAT:

```
add     sub     mul     div
rem     and     or      xor
abs     neg     min     max
clz     ctz     floor   nearest
```

The following list of mnemonics all use different names to their WAT equivalents, however they are otherwise the same as the mnemonics in the previous list:

```
ceiling     ; ceil
lop         ; trunc
root        ; sqrt
sign        ; copysign
nsa         ; popcnt
```

Below are a handful of examples, showing how the grammar for arithmetic and logic looks in practice:

```
add i64
abs f64
xor i32
div u32
rem s64
```

The Shift & Rotate Mnemonics

PHANTASM uses the `shift` and `rotate` mnemonics for all of the `shl`, `shr`, `rotl` and `rotr` WAT instructions (including gnostic instructions, like `i32.shr_s` *et cetera*).

The grammar is similar to the grammar for arithmetic and logic instructions, but appends a `left` or `right` qualifier to the end:

```
mnemonic numtype carousel-qualifier
```

Naturally, the numtype must be either `i32` or `i64`, except for shift-right-instructions, where the type is gnostic (`u32`, `s32`, `u64` or `s64`). For example:

```
shift i32 left
shift s32 right
rotate i32 right
rotate i64 left
shift u64 right
```

The Call & Invoke Mnemonics

PHANTASM copies the `call` mnemonic from WAT, but uses the `invoke` mnemonic for `call_indirect`.

The grammar for the call-instruction simply appends an identity to the mnemonic:

call identity

The grammar for the invoke-instruction requires a type, and an optional identity, prefixed by the `via` keyword:

invoke type [**via** identity]

The type can be expressed with a type-reference or type-expression, and the identity specifies the table, defaulting to zero. For example:

```
call 12
call $helper
invoke type $binop
invoke i32, i32 → i32
invoke type $check via $checks
invoke i32 → void via 2
```

The Branch Instructions

PHANTASM renames all three of the WAT branch-instructions, from `br`, `br_if` and `br_table` to `jump`, `fork` and `exit`, respectively.

In PHANTASM terms, a *jump* is unconditional, while a *fork* is a conditional jump. An *exit* always exits the current block, but uses an operand and an (immediate) array of indices to determine which exit to take.

The jump-instruction and fork-instruction use a grammar that appends an identity (for the target block) to the mnemonic:

mnemonic identity

The exit-instruction requires one or more (space-separated) identities (with the same order as WAT):

```
mnemonic identity+
```

For example:

```
jump 0
fork $block
exit 1 0 2 $loop
```

The Block Instructions

PHANTASM copies the `block` and `loop` mnemonics from WAT, and also uses `else`. However, PHANTASM uses the `branch` mnemonic instead of `if` (so `else` blocks optionally follow `branch` blocks).

All three branch instructions (`block`, `loop` and `branch`) use the same grammar, which reuses the function grammar and the `of` keyword:

```
mnemonic [identifier] of type block
```

When present, the identifier is bound to the implied label, and (like the invoke-instruction) the type can be expressed with a type-reference or a type-expression. For example:

```
loop $loop of type 2

    ; instructions ...

block of i32, i32 → void

    ; instructions ...

branch $branch of type $check

    ; instructions ...

else

    ; instructions ...
```

Note: The blocks must be indented (block-instructions cannot contain inline-blocks, and inline-blocks cannot contain block-instructions).

Note: PHANTASM does not require whitelines anywhere in the grammar, but you are recommended to use (exactly) one before and after each block header, block, directive and docstring, as well as between chunks of related statements, directives and commands. Adjacent whitelines are never recommended.

The Block Instruction Sugar

Block-instructions (and only block-instructions) support sugar for expressing their type, using the `with` keyword to prefix a valtype:

```
mnemonic [identifier] with valtype block
```

The above grammar is equivalent to the following grammar:

```
mnemonic [identifier] of valtype → void block
```

For example, the following two headers are equivalent:

```
block $block with i32
block $block of i32 → void
```

Likewise, the following two headers are equivalent:

```
loop with void
loop of void → void
```

The Select Mnemonic

PHANTASM copies the `select` mnemonic from WAT, with the same grammar:

```
select [valtype]
```

For example:

```
select
select i32
select pointer
```

The Memory & Table Instructions

PHANTASM uses the `grow`, `size` and `fill` mnemonics from WAT, but the component type (`memory` or `table`) follows the mnemonic, with an optional identity for the memory or table:

```
mnemonic component-reference [identity]
```

For example:

```
grow memory
size memory 1
fill table $opcodes
```

The Drop Mnemonic

PHANTASM uses the `drop` mnemonic from WAT, with the same grammar as the memory and table instructions (above), except that the component-reference must reference a bank (and therefore, the identity is required):

```
mnemonic component-reference identity
```

For example:

```
drop memory bank 0
drop table bank $extensions
```

The Copy & Initialization Instructions

PHANTASM uses the `copy` mnemonic for the WAT instructions `copy` and `init`, with the following grammar:

```
mnemonic component-reference [to identity]
```

The component-reference defines the *location* (the component to copy from). It can reference a memory or table (implementing the WAT copy-instruction), or a memory bank or table bank (implementing the WAT init-instruction).

The optional identity (prefixed by the `to` keyword) is used to identify the *destination* (the component to copy to). It defaults to zero. Given that banks are readonly, the destination type can be inferred from the location type (`memory` for memories and memory banks, and `table` for tables and table banks).

```
copy table 1
copy memory to 5
copy table $opcodes
copy table $messages to $messages

copy table bank $extensions to $opcodes
copy memory bank 0 to 5
copy table bank 1
```

The Is & Not Mnemonics

PHANTASM uses the `is` and `not` mnemonics for the WAT mnemonics `eqz`, `eq`, `ne`, `gt`, `lt`, `ge` and `le`, using the following grammar:

```
mnemonic test
```

When the mnemonic is `is`, the test can be one of `null`, `zero`, `equal`, `less` or `more`. When the mnemonic is `not`, the test can only be one of `equal`, `less` or `more`. For example:

```
is zero
is null
not more
is less
not equal
```

The Conversion Instructions

PHANTASM reuses the `wrap`, `convert`, `promote` and `demote` mnemonics with the following grammar:

```
mnemonic numtype to numtype
```

For example:

```
wrap i64 to i32
convert s32 to f64
promote f32 to f64
demote f64 to f32
```

PHANTASM also uses the same grammar for its `cast` and `lop` mnemonics, which replaces the WAT mnemonics `reinterpret` and `trunc`, respectively. For example:

```
cast i32 to f32
cast i64 to f64
cast f32 to i32

lop f32 to u32
lop f32 to s64
lop f64 to u32
```

PHANTASM uses a very similar grammar for its version of the WAT `trunc_sat` instructions, which replace the `to` keyword with the `sop` keyword:

```
lop f32 sop u32
lop f32 sop s64
lop f64 sop u32
```

The PHANTASM `extend` mnemonic is used for the two WAT extend-instructions that sign-extend and zero-extend the operand, producing a 64-bit result from a 32-bit operand. For example:

```
extend u32 to i64
extend s32 to i64
```

Note: The other WAT extend-instructions use the PHANTASM `expand` mnemonic (see below).

The Expand Instructions

PHANTASM uses the `expand` mnemonic for those WAT extend-instructions that sign-extend the least-significant bytes of the operand, mutating the most significant bytes to produce a result that is the same width (as the operand).

The expand-instructions reuse the grammar of the `load` and `store` instructions (which is also reused by most of the atomic instructions, described below).

```
mnemonic numtype as subtype
```

The numtype describes the type of the operand (and result), while the subtype describes the smaller value (that will be extended).

```
expand i32 as s8
expand i32 as s16
expand i64 as s32
```

The Atomic Instructions

Most of the PHANTASM atomic-instructions (`load`, `store`, `add`, `sub`, `and`, `or`, `xor`, `swap` and `broker`) use the same grammar as the non-atomic `load` and `store` instructions, but prefixed by the `atomic` qualifier:

```
atomic mnemonic numtype [as datatype] [in identity] [at number-literal]
```

Note: In practice, the numtype of any atomic-operation will be an integer.

PHANTASM uses the `swap` and `broker` mnemonics for the WAT mnemonics `xchg` and `cmpxchg`.

The various constructs of the grammar are interpreted in the same way as the `load` and `store` instructions. For example:

```
atomic load i32
atomic store i64 as i8
atomic add i32 in $memory
atomic broker i32 as i16
atomic xor i64 as i32 in $RAM at 2
```

The atomic-wait-instructions use a similar grammar to the regular atomic-instructions, just without the (optional) datatype:

```
atomic wait numtype [in identity] [at number-literal]
```

The atomic-notify-instruction uses the same grammar as the atomic-wait instructions, but without the (required) numtype:

```
atomic notify [in identity] [at number-literal]
```

Below are a few examples of how the `wait` and `notify` mnemonics are used:

```
atomic notify
atomic wait i32
atomic wait i64 at 4
atomic notify in $memory at 2
```

The atomic-fence-instruction is simply written as follows:

atomic fence

The atomic-fence instruction completes the section on PHANTASM instructions, and with it, the section on function-definitions. The rest of this document describes the component-definitions for types, memories and tables (including banks).

Memory Definitions

Memory-definitions append a required initializer, known as a *memory-primer*, to the grammar for memory-specifiers:

[**shared**] **memory** [identifier] limits [primer]

When omitted, the primer is implicitly empty. When present, the primer is a block (which can be inlined), and must contain one or more *segments*. Segments are defined using a *segment-directive*, followed by one or more *memory-commands*, and are compiled to active (memory) segments (within the binary).

Segment Directives

A segment-directive simply specifies an offset, using the following grammar:

segment block

The block is a constant expression (grammatically, just a block of instructions) that defines the offset, and it can be inlined. For example:

```
segment thus push #100
```

Segment-directives also support some syntactic sugar, using the `at` keyword with the following grammar:

segment at number-literal

The sugar compiles to a push-instruction, using the number literal as its immediate, so for example, the following code is equivalent to the previous example:

```
segment at #100
```

Memory Commands

Memory-commands define the chunks of data that are used to populate memories when they are initialized. All memory-commands use the following grammar:

```
mnemonic value
```

The mnemonic must be one of `i8`, `i16`, `i32`, `i64` or `utf8`. When the mnemonic is a numtype, the value must be a valid number literal (for the given type). When the mnemonic is `utf8`, the value must be a string literal.

As you might expect, the mnemonic can be omitted and inferred from the previous command, when they use the same type, and multiple commands can be grouped on the same line. For example:

```
utf8 "Hello, World!"
i8 #10, #20, i16 #30, #40
```

Memory Primers

Memory-primers (when present) are blocks that contain one or more segments, each defined by a segment-directive, followed by one or more memory-commands. For example:

```
define shared memory $memory with 16 to 256

    segment at #500

    i8 #10, #20, #30, #40
    i16 #10, #20, #30, #40

    segment at #1000

    utf8 "Hello, World!"
    i8 #10, #20, i16 #30, #40
```

The grammar permits the first segment in a primer to omit its segment-directive (and the offset is implicitly zero). For example:

```
define memory with 1 max

    utf8 "Congratulations ... You Won!"
    utf8 "Commiserations ... You Lost!"

    segment at 64

    i8 #10, #20, #30, #40
    i16 #10, #20, #30, #40
```

When the segment-directive is omitted from a primer with only one segment, the primer can be inlined (inline-blocks cannot contain directives, as directives must occupy their own logical line within a block). For example:

```
define memory with 1 max thus i8 #10, #20, #30
```

Memory Bank Definitions

Memory banks are compiled to passive (memory) segments (within the binary), serving as little ROMs that can be copied from at runtime.

PHANTASM compiles all passive segments ahead of any active segments, so memory banks and table banks each (effectively) have their own indexspaces.

Note: The indices of active segments are not accessible (nor useful) to the user. Attempting to access one (using a number literal to specify its index) will result in a compile-time error.

The grammar for memory-bank-definitions is derived from the grammar for memory-definitions. The limits-construct is removed, while the primer is required:

```
memory bank [identifier] primer
```

Note: Bank primers cannot contain segment-directives (a bank *is* a segment).

Memory-bank-definitions are used by define-statements. For example:

```
define memory bank $bank

    utf8 "Hello, World!"
    i8 #10, #20, i16 #30, #40
```

As ever, blocks that only contain a single line of commands can be inlined:

```
define memory bank thus i8 #10, #20, i16 #30, #40
```

Table Definitions

Table-definitions extend the grammar of table-specifiers, appending a optional block that defines a *table-primer*:

```
table-qualifier table [identifier] limits primer
```

The various constructs of the above grammar have been described already (in the Table Specifiers and Memory Definitions sections): The qualifier is one of `pointer`, `proxy` or `mixed`, and the limits are defined slot-wise. However, in a table-definition, the primer contains table-commands (instead of memory-commands), and table-definitions can only include primers if the table type is `pointer`.

Table Commands

Table-commands use the same grammar as memory-commands:

```
mnemonic value
```

The mnemonic must (currently) be `pointer`. More table-commands will be added soon. The value is a reference to a function that can be expressed as an identity or `null`. As usual, the mnemonic can be inferred from the previous command. For example:

```
pointer 14, 9, null, $helper, null
```

Table Primers

Table-primers (when present) are blocks that contain one or more (table) segments, each defined by a segment-directive (which can be omitted, and implicitly zero, for the first segment), followed by one or more table-commands. For example:

```
define pointer table $opcodes with #1000 to #2000

    pointer $nop, $jsr, $rts

    segment at #100

    pointer 1, 2, 3, null, null, null
    pointer $foo, $bar, $spam, $eggs

define pointer table with 256 max thus pointer $nop, $jsr, $rts
```

Table Bank Definitions

Table banks are compiled to passive (table) segments (within the binary), just like memory banks.

The grammar for table-bank-definitions is the same as for memory banks, except the component name is replaced with the table type:

```
type bank [identifier] primer
```

Currently (due to the WebAssembly Specification), table banks must use the `pointer` type, so in practice, the effective grammar is more restrictive than above:

```
pointer bank [identifier] primer
```

Bank primers cannot contain segment-directives, so the primer will consist of one or more table-commands, which must each use the `pointer` mnemonic (at the moment).

Table-bank-definitions are used by define-statements. For example:

```
define pointer bank $bank thus pointer $nop, $jsr, $rts
```

Type Definitions

Type-definitions define function types, using a simple grammar:

```
type [identifier] as type-expression
```

Type-definitions are only used in define-statements. For example:

```
define type $binop as i32, i32 → i32
```

When the user defines a type explicitly (as above), the type is known as an *explicit type*.

Every type-definition (within a module) must be unique (the user cannot duplicate type-definitions).

There are six places in the grammar where a type can be referenced (function-specifiers, function-definitions, the three block-instructions and the invoke-instruction). In all six cases, the type can either be referenced (with a type-reference, like `type $binop`) or expressed (with a type-expression, like `i32, i32 -> i32`).

When a type-reference is used, it must reference an explicit type. When a type-expression is used, it may express any type.

When a type-expression expresses a unique type, the type is defined implicitly, and is known as an *implicit type*. If the type has been defined already (implicitly or explicitly), the type-expression compiles to a reference to the same type.

As with banks (where every passive segment is compiled ahead of any active segments), explicit types are compiled before any implicit types (and the indices of implicit types are inaccessible to the user), effectively giving the explicit types their own indexspace.

Note: Every component type (register, function, memory, table, memory bank, table bank and function type) has its own indexspace.

Note: As explicit type-definitions must define unique types, there is no way for a module to define a duplicate type.

The lack of duplicate types neatly sidesteps the issue WAT has, where a function-definition cannot reference a type *and* bind identifiers to the parameters of the type, without duplicating the type or requiring the author to include both a type-reference *and* a signature (which, naturally, must express the same type). In PHANTASM, the author would just use a signature as normal, without any special implications.

Going Forwards

Now that you have finished reading the Crashcourse, you should grab (and maybe print) a copy of the _PHANTASM: Abstract Grammar Cheatsheet_, which summarizes the syntax using a simple, BNF-style pseudocode.

The project wiki also contains a growing collection of useful articles:

• _Installation_: How to get the code and use it.
• _String Literals_: Details string literals, particularly interpolation.
• _Number Literals_: Details number literals, particularly exponentiation and evaluation.
• _Roadmap & Status_: Details number literals, particularly exponentiation and evaluation.

Please feel free to use the project issue tracker as a general forum for any relevant discussions.

Welcome to the PHANTASM community!